


Branch: master ▾

[ML](#) / Bai tap 2.ipynb

Find file

Copy path

 **tranquoclongt1** full HW1.ex2

8511cc9 2 days ago

1 contributor

1.55 MB

Tran Quoc Long - 14520490

Bài tập 2: Handwritting digits - clustering

K-means


```
In [19]: #import libs
from time import time
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
In [20]: #import scikit-learn
from sklearn import metrics
from sklearn.cluster import KMeans
from sklearn.datasets import load_digits
```

```
In [21]: digits = load_digits();
print(digits.data.shape);
```

```
(1797, 64)
```

```
In [22]: %matplotlib inline
plt.gray();
plt.matshow(digits.images[0]);
```



```
In [23]: nClusters = 10
model1 = KMeans(nClusters)
labels_kmeans = model1.fit_predict(digits.data)
```


```
In [24]: df = pd.DataFrame({'labels':labels_kmeans,'Truth labels':digits.target})
ct = pd.crosstab(df['labels'],df['Truth labels'])
print(ct)
```

Truth labels	0	1	2	3	4	5	6	7	8	9
labels										
0	0	1	13	155	0	0	0	0	2	6
1	0	0	2	13	0	43	0	0	52	139
2	0	2	0	0	0	1	177	0	2	0
3	0	55	2	0	5	0	1	2	6	20
4	0	24	148	0	0	0	0	0	3	0
5	1	0	0	0	164	1	0	0	0	0
6	0	99	8	7	3	0	2	2	99	1
7	177	0	1	0	0	0	1	0	0	0
8	0	0	3	6	9	0	0	170	2	7
9	0	1	0	2	0	137	0	5	8	7

```
In [25]: n = 10
%matplotlib inline
plt.matshow(digits.images[n])
print('Predict Label:', labels_kmeans[n])

print('Truth: ', digits.target[n])
```

```
Predict Label: 7
Truth: 0
```




Visualization - Kmeans


```
In [26]: #import libs
import numpy as np
from sklearn.decomposition import PCA
```

PCA

```
In [27]: nComponents = 2
vPCA = PCA(nComponents)
digitData_to_2D = vPCA.fit_transform(digits.data)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_kmeans, s=20)
plt.show()
```



```
In [29]: plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= digits.target, s=20)
plt.show()
```



In []:

Spectral clustering

```
In [44]: # Spectral_clustering

from sklearn.cluster import spectral_clustering
from sklearn.feature_extraction import image
import numpy as np
```

```

from sklearn.neighbors import DistanceMetric
from sklearn.metrics.pairwise import cosine_similarity

# dist = DistanceMetric.get_metric('euclidean')
# graph=dist.pairwise(digits.data)

graph = cosine_similarity(digits.data)
label_spectral = spectral_clustering(graph, n_clusters=10)

```

```

In [45]: df1 = pd.DataFrame({'labels':label_spectral,'Truth labels':digits.target})
ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
print(ct2)

```


Truth labels	0	1	2	3	4	5	6	7	8	9
labels										
0	0	58	5	5	1	0	0	15	40	36
1	0	0	1	146	0	0	0	0	6	2
2	0	0	0	4	0	157	0	0	3	3
3	0	2	0	1	0	2	172	0	13	0
4	0	0	0	16	0	20	2	0	7	133
5	177	0	1	0	1	1	0	0	0	3
6	0	0	2	2	11	0	0	154	3	2
7	0	36	115	4	0	0	0	0	1	0
8	1	0	0	0	163	2	0	0	0	0
9	0	86	53	5	5	0	7	10	101	1

```

In [46]: n = 15
plt.matshow(digits.images[n])
print('lables_predict:',label_spectral[n])
print(' True: ', digits.target[n])

```

lables_predict: 2
True: 5




Visualization - Spectral Clustering

```

In [15]: plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= label_spectral, s=20)
plt.show()

```



Visualize results to compare - Using PCA

```

In [58]: fig = plt.figure(figsize=(15,4))
fig.suptitle('Comparison results of methods', fontsize=20)

ax = fig.add_subplot(1,3,1)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= label_spectral, s=20)

```

```


plt.figure(figsize=(10,10))
ax.set_title('Spectral Clustering')

ax = fig.add_subplot(1,3,2)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_kmeans, s=20)
ax.set_title('Kmeans Clustering')

ax = fig.add_subplot(1,3,3)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= digits.target, s=20)
ax.set_title('Truth Result')

```

Out[58]: <matplotlib.text.Text at 0x23de36f9be0>



DBSCAN

```

In [26]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import DBSCAN
from sklearn import metrics
from sklearn.datasets.samples_generator import make_blobs
from sklearn.preprocessing import StandardScaler

#import scikit-Learn
from sklearn import metrics
from sklearn.cluster import KMeans
from sklearn.datasets import load_digits
from sklearn.preprocessing import scale
from sklearn.decomposition import PCA

```

```

In [27]: digits = load_digits()
data = digits.data
data = StandardScaler().fit_transform(data)

n_samples, n_features = data.shape
n_digits = len(np.unique(digits.target))
labels = digits.target

sample_size = 300

print("n_digits: %d, \t n_samples %d, \t n_features %d"
      % (n_digits, n_samples, n_features))

n_digits: 10,      n_samples 1797,      n_features 64

```

```

In [28]: print(data)

[[ 0.         -0.33501649 -0.04308102 ..., -1.14664746 -0.5056698
  -0.19600752]
 [ 0.         -0.33501649 -1.09493684 ...,  0.54856067 -0.5056698
  -0.19600752]
 [ 0.         -0.33501649 -1.09493684 ...,  1.56568555  1.6951369
  -0.19600752]
 ...,
 [ 0.         -0.33501649 -0.88456568 ..., -0.12952258 -0.5056698
  -0.19600752]
 [ 0.         -0.33501649 -0.67419451 ...,  0.8876023  -0.5056698
  -0.19600752]
 [ 0.         -0.33501649  1.00877481 ...,  0.8876023  -0.26113572
  -0.19600752]]

```

```

In [45]: db = DBSCAN(eps=1, min_samples=1, algorithm='kd_tree').fit(data)

```

In [46]: `print(db)`

```
DBSCAN(algorithm='kd_tree', eps=1, leaf_size=30, metric='euclidean',
        metric_params=None, min_samples=1, n_jobs=1, p=None)
```

In [47]: `print(db.labels_)`
`print(sum(db.labels_== -1))`


```
[ 0  1  2 ..., 1793 1794 1795]
0
```

In [48]: `import pandas as pd`
`df1 = pd.DataFrame({'labels':db.labels_, 'Truth labels':digits.target})`
`ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])`
`print(ct2)`

Truth labels	0	1	2	3	4	5	6	7	8	9
labels										
0	1	0	0	0	0	0	0	0	0	0
1	0	1	0	0	0	0	0	0	0	0
2	0	0	1	0	0	0	0	0	0	0
3	0	0	0	1	0	0	0	0	0	0
4	0	0	0	0	1	0	0	0	0	0
5	0	0	0	0	0	1	0	0	0	0
6	0	0	0	0	0	0	1	0	0	0
7	0	0	0	0	0	0	0	1	0	0
8	0	0	0	0	0	0	0	0	1	0
9	0	0	0	0	0	0	0	0	0	1
10	1	0	0	0	0	0	0	0	0	0
11	0	1	0	0	0	0	0	0	0	0
12	0	0	1	0	0	0	0	0	0	0
13	0	0	0	1	0	0	0	0	0	0
14	0	0	0	0	1	0	0	0	0	0
15	0	0	0	0	0	1	0	0	0	0
16	0	0	0	0	0	0	1	0	0	0
17	0	0	0	0	0	0	0	1	0	0
18	0	0	0	0	0	0	0	0	1	0
19	0	0	0	0	0	0	0	0	0	1
20	1	0	0	0	0	0	0	0	0	0
21	0	1	0	0	0	0	0	0	0	0
22	0	0	1	0	0	0	0	0	0	0
23	0	0	0	1	0	0	0	0	0	0
24	0	0	0	0	1	0	0	0	0	0
25	0	0	0	0	0	1	0	0	0	0
26	0	0	0	0	0	0	1	0	0	0
27	0	0	0	0	0	0	0	1	0	0
28	0	0	0	0	0	0	0	0	1	0
29	0	0	0	0	0	0	0	0	0	1
...
1766	0	0	0	0	1	0	0	0	0	0
1767	1	0	0	0	0	0	0	0	0	0
1768	0	0	0	0	0	1	0	0	0	0
1769	0	0	0	1	0	0	0	0	0	0
1770	0	0	0	0	0	0	1	0	0	0
1771	0	0	0	0	0	0	0	0	0	1
1772	0	0	0	0	0	0	1	0	0	0
1773	0	1	0	0	0	0	0	0	0	0
1774	0	0	0	0	0	0	0	1	0	0
1775	0	0	0	0	0	1	0	0	0	0
1776	0	0	0	0	1	0	0	0	0	0
1777	0	0	0	0	1	0	0	0	0	0
1778	0	0	0	0	0	0	0	1	0	0
1779	0	0	1	0	0	0	0	0	0	0
1780	0	0	0	0	0	0	0	0	1	0
1781	0	0	1	0	0	0	0	0	0	0
1782	0	0	1	0	0	0	0	0	0	0
1783	0	0	0	0	0	1	0	0	0	0
1784	0	0	0	0	0	0	0	1	0	0
1785	0	0	0	0	0	0	0	0	0	1
1786	0	0	0	0	0	1	0	0	0	0
1787	0	0	0	0	1	0	0	0	0	0
1788	0	0	0	0	0	0	0	0	1	0
1789	0	0	0	0	0	0	0	0	1	0
1790	0	0	0	0	1	0	0	0	0	0
1791	0	0	0	0	0	0	0	0	0	1
1792	1	0	0	0	0	0	0	0	0	0
1793	0	0	0	0	0	0	0	0	1	0
1794	0	0	0	0	0	0	0	0	0	1
1795	0	0	0	0	0	0	0	0	1	0

[1796 rows x 10 columns]

```
In [49]: nComponents = 2
vPCA = PCA(nComponents)
digitData_to_2D = vPCA.fit_transform(digits.data)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= db.labels_, s=20)
plt.show()
```



Agglomerative Clustering

```
In [53]: from sklearn.cluster import AgglomerativeClustering
```

```
In [57]: Agglomerative_model = AgglomerativeClustering(n_clusters = 10)
```

```
In [58]: db = Agglomerative_model.fit(data)
```


```
In [59]: print(db.labels_)
```

```
[5 1 1 ..., 1 1 1]
```

```
In [60]: import pandas as pd
df1 = pd.DataFrame({'labels':db.labels_, 'Truth labels':digits.target})
ct2 = pd.crosstab(df1['labels'], df1['Truth labels'])
print(ct2)
```

Truth labels labels	0	1	2	3	4	5	6	7	8	9
0	0	1	0	0	1	168	0	1	1	3
1	0	150	15	11	4	0	1	1	168	38
2	0	0	1	0	1	0	0	0	0	0
3	0	27	160	4	0	1	0	0	3	0
4	0	0	1	168	0	12	0	1	2	135
5	178	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	1	180	0	0	1
7	0	0	0	0	12	0	0	25	0	3
8	0	4	0	0	163	0	0	0	0	0
9	0	0	0	0	0	0	0	151	0	0

```
In [61]: plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= db.labels_, s=20)
plt.show()
```



Comparison of cluster methods

```
In [84]: #import libs
```


```
from time import time
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
In [85]: #import scikit-learn
from sklearn import metrics
from sklearn.cluster import KMeans, spectral_clustering, DBSCAN, AgglomerativeClustering
from sklearn.datasets import load_digits
from sklearn.neighbors import DistanceMetric
from sklearn.metrics.pairwise import cosine_similarity
```

```
In [86]: digits = load_digits();
print(digits.data.shape);
```

```
(1797, 64)
```

```
In [87]: %matplotlib inline
plt.gray();
plt.matshow(digits.images[0]);
```



Clustering

```
In [100]: #Kmeans
nClusters = 10
t0 = time()
kmeans_model = KMeans(nClusters)
t_kmeans = time() - t0
labels_kmeans = kmeans_model.fit_predict(digits.data)
```

```
In [101]: #Spectral_clustering
t0 = time()
graph = cosine_similarity(digits.data)
t_spectral = time() - t0
labels_spectral = spectral_clustering(graph, n_clusters=10)
```

```
In [102]: #DBSCAN
data = digits.data
t0 = time()
data = StandardScaler().fit_transform(data)
labels_dbscan = DBSCAN(eps=1, min_samples=1, algorithm='kd_tree').fit_predict(data)
t_dbscan = time() - t0
```

```
In [103]: #Agglomerative Clustering
t0 = time()
Agglomerative_model = AgglomerativeClustering(n_clusters = nClusters)
labels_AgglomerativeClustering = Agglomerative_model.fit_predict(data)
t_agg = time() - t0
```

Cross table

```
In [104]: #Kmeans
print('Kmeans:\n')
df1 = pd.DataFrame({'labels':labels_kmeans,'Truth labels':digits.target})
ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
print(ct2)
```

```
#Spectral_clustering
```



```
#Spectral clustering
print('\n\n')
print('Spectral clustering:\n')
df1 = pd.DataFrame({'labels':labels_spectral,'Truth labels':digits.target})
ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
print(ct2)

#DBSCAN
print('\n\n')
print('DBSCAN:\n')
df1 = pd.DataFrame({'labels':labels_dbscan,'Truth labels':digits.target})
ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
print(ct2)

#Agglomerative Clustering
print('\n\n')
print('Agglomerative Clustering:\n')
df1 = pd.DataFrame({'labels':labels_AgglomerativeClustering,'Truth labels':digits.target})
ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
print(ct2)
```

Kmeans:

Truth labels labels	0	1	2	3	4	5	6	7	8	9
0	0	99	8	7	3	0	2	2	102	2
1	0	0	2	11	0	42	0	0	48	139
2	0	0	3	7	9	0	0	175	5	7
3	0	1	13	155	0	1	0	0	4	6
4	0	55	2	0	5	0	1	2	6	20
5	0	2	0	0	0	1	177	0	2	0
6	177	0	1	0	0	0	1	0	0	0
7	1	0	0	0	164	2	0	0	0	0
8	0	1	0	2	0	136	0	0	4	6
9	0	24	148	1	0	0	0	0	3	0

Spectral clustering:

Truth labels labels	0	1	2	3	4	5	6	7	8	9
0	177	0	1	0	1	1	0	0	0	3
1	1	0	0	0	163	2	0	0	0	0
2	0	0	1	145	0	0	0	0	6	2
3	0	58	5	5	1	0	0	14	36	35
4	0	0	0	4	0	157	0	0	3	3
5	0	0	2	2	11	0	0	155	3	2
6	0	0	0	16	0	20	3	0	7	134
7	0	2	0	1	0	2	172	0	14	0
8	0	36	115	4	0	0	0	0	1	0
9	0	86	53	6	5	0	6	10	104	1

DBSCAN:

Truth labels labels	0	1	2	3	4	5	6	7	8	9
0	1	0	0	0	0	0	0	0	0	0
1	0	1	0	0	0	0	0	0	0	0
2	0	0	1	0	0	0	0	0	0	0
3	0	0	0	1	0	0	0	0	0	0
4	0	0	0	0	1	0	0	0	0	0
5	0	0	0	0	0	1	0	0	0	0
6	0	0	0	0	0	0	1	0	0	0
7	0	0	0	0	0	0	0	1	0	0
8	0	0	0	0	0	0	0	0	1	0
9	0	0	0	0	0	0	0	0	0	1
10	1	0	0	0	0	0	0	0	0	0
11	0	1	0	0	0	0	0	0	0	0
12	0	0	1	0	0	0	0	0	0	0
13	0	0	0	1	0	0	0	0	0	0
14	0	0	0	0	1	0	0	0	0	0
15	0	0	0	0	0	1	0	0	0	0
16	0	0	0	0	0	0	1	0	0	0
17	0	0	0	0	0	0	0	1	0	0
18	0	0	0	0	0	0	0	0	1	0

```

19      0 0 0 0 0 0 0 0 0 0 1
20      1 0 0 0 0 0 0 0 0 0 0
21      0 1 0 0 0 0 0 0 0 0 0
22      0 0 1 0 0 0 0 0 0 0 0
23      0 0 0 1 0 0 0 0 0 0 0
24      0 0 0 0 1 0 0 0 0 0 0
25      0 0 0 0 0 1 0 0 0 0 0
26      0 0 0 0 0 0 1 0 0 0 0
27      0 0 0 0 0 0 0 1 0 0 0
28      0 0 0 0 0 0 0 0 1 0 0
29      0 0 0 0 0 0 0 0 0 1 0
...
1766    0 0 0 0 1 0 0 0 0 0 0
1767    1 0 0 0 0 0 0 0 0 0 0
1768    0 0 0 0 0 1 0 0 0 0 0
1769    0 0 0 1 0 0 0 0 0 0 0
1770    0 0 0 0 0 0 1 0 0 0 0
1771    0 0 0 0 0 0 0 0 0 0 1
1772    0 0 0 0 0 0 1 0 0 0 0
1773    0 1 0 0 0 0 0 0 0 0 0
1774    0 0 0 0 0 0 0 1 0 0 0
1775    0 0 0 0 0 1 0 0 0 0 0
1776    0 0 0 0 1 0 0 0 0 0 0
1777    0 0 0 0 1 0 0 0 0 0 0
1778    0 0 0 0 0 0 0 1 0 0 0
1779    0 0 1 0 0 0 0 0 0 0 0
1780    0 0 0 0 0 0 0 0 0 1 0
1781    0 0 1 0 0 0 0 0 0 0 0
1782    0 0 1 0 0 0 0 0 0 0 0
1783    0 0 0 0 0 1 0 0 0 0 0
1784    0 0 0 0 0 0 0 1 0 0 0
1785    0 0 0 0 0 0 0 0 0 1 0
1786    0 0 0 0 0 1 0 0 0 0 0
1787    0 0 0 0 1 0 0 0 0 0 0
1788    0 0 0 0 0 0 0 0 0 1 0
1789    0 0 0 0 0 0 0 0 0 1 0
1790    0 0 0 0 1 0 0 0 0 0 0
1791    0 0 0 0 0 0 0 0 0 0 1
1792    1 0 0 0 0 0 0 0 0 0 0
1793    0 0 0 0 0 0 0 0 0 1 0
1794    0 0 0 0 0 0 0 0 0 0 1
1795    0 0 0 0 0 0 0 0 0 1 0

```

[1796 rows x 10 columns]

Agglomerative Clustering:

Truth labels	0	1	2	3	4	5	6	7	8	9
0	0	1	0	0	1	168	0	1	1	3
1	0	150	15	11	4	0	1	1	168	38
2	0	0	1	0	1	0	0	0	0	0
3	0	27	160	4	0	1	0	0	3	0
4	0	0	1	168	0	12	0	1	2	135
5	178	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	1	180	0	0	1
7	0	0	0	0	12	0	0	25	0	3
8	0	4	0	0	163	0	0	0	0	0
9	0	0	0	0	0	0	0	151	0	0

Comparison

```

In [112]: #print frame
print("n_digits: %d, \t n_samples %d, \t n_features %d"
      % (n_digits, n_samples, n_features))

print(82 * '_')
print('init\ttime\tinertia\tthomo\tcompl\tv-meas\tARI\tAMI\tsilhouette')

data = digits.data
#define a function to measure and print out
def bench_clustering(method_name, time_, labels):
    print('%-9s\t%.2fs\t%.3f\t%.3f\t%.3f\t%.3f\t%.3f\t%.3f'
          % (method_name, time_,

```

```

metrics.homogeneity_score(digits.target, labels),
metrics.completeness_score(digits.target, labels),
metrics.v_measure_score(digits.target, labels),
metrics.adjusted_rand_score(digits.target, labels),
metrics.adjusted_mutual_info_score(digits.target, labels),
metrics.silhouette_score(data, labels,
                           metric='euclidean',
                           sample_size=sample_size)))

#Kmeans
bench_clustering('K-means', t_kmeans, labels_kmeans)
#Spectral_clustering
bench_clustering('spectral', t_spectral, labels_spectral)

#Agglomerative clustering
bench_clustering('Agg.', t_agg, labels_AgglomerativeClustering)
#DBSCAN ==> Problems with raw data
#bench_clustering('DBSCAN', t_dbscan, labels_dbscan)
print('-----\nProblems with raw data cause noise with DBSCAN method')

```

n_digits: 10, n_samples 1797, n_features 64

init	time	inertia	homo	compl	v-meas	ARI	AMI	silhouette
K-means	0.00s	0.739	0.747	0.743	0.669	0.736	0.178	
spectral	0.03s	0.711	0.716	0.713	0.625	0.708	0.165	
Agg.	0.15s	0.758	0.836	0.796	0.664	0.756	0.127	

Problems with raw data cause noise with DBSCAN method

Visualization

```

In [149]: nComponents = 2
vPCA = PCA(nComponents)
digitData_to_2D = vPCA.fit_transform(digits.data)

fig = plt.figure(figsize=(15,15))
fig.suptitle('Comparison results of methods', fontsize=20)

ax = fig.add_subplot(3,2,1)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= label_spectral, s=20)
ax.set_title('Spectral Clustering')

ax = fig.add_subplot(3,2,2)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_dbscan, s=20)
ax.set_title('DBSCAN Clustering')

ax = fig.add_subplot(3,2,3)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_kmeans, s=20)
ax.set_title('K-means')

ax = fig.add_subplot(3,2,4)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_AgglomerativeClustering, s=20)
ax.set_title('Agglomerative clustering')

ax = fig.add_subplot(3,1,3)
plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= digits.target, s=20)
ax.set_title('Target Result')

```

Out[149]: <matplotlib.text.Text at 0x22075511a90>

Comparison results of methods

