

## Bài tập 2: Handwritten digits - clustering

```
1 Thực hiện các phép cluster trên bộ dữ liệu handwritten digits
2 Nội dung bao gồm trong file:
3 1. Chạy thử các hàm cluster và các hàm liên quan
4     - Kmeans
5     - Spectral clustering
6     - DBSCAN
7     - Agglomerative clustering
8     - Cross table
9     - Figure to visualize result
10    - Show centroid of Kmeans
11 2. Nội dung thực hành 2
```

### Chạy thử các hàm cluster và các hàm liên quan

#### K-means


```
In [1]: 1 #import libs
        2 from time import time
        3 import numpy as np
        4 import matplotlib.pyplot as plt
        5 import pandas as pd
```

```
In [2]: 1 #import scikit-learn
        2 from sklearn import metrics
        3 from sklearn.cluster import KMeans
        4 from sklearn.datasets import load_digits
```

```
In [3]: 1 digits = load_digits();
        2 print(digits.data.shape);
```

(1797, 64)

```
In [4]: 1 %matplotlib inline
        2 #plt.gray();
        3 plt.matshow(digits.images[0]);
```



```
In [5]: 1 nClusters = 10
        2 model1 = KMeans(nClusters)
        3 labels_kmeans = model1.fit_predict(digits.data)
```


```
In [6]: 1 df = pd.DataFrame({'labels':labels_kmeans,'Truth labels':digits.target})
        2 ct = pd.crosstab(df['labels'],df['Truth labels'])
        3 print(ct)
```

Truth labels	0	1	2	3	4	5	6	7	8	9
0	0	0	3	7	9	0	0	175	5	7
1	0	1	0	2	0	136	0	0	4	6
2	177	0	1	0	0	0	1	0	0	0
3	0	0	2	9	0	42	0	0	48	138
4	0	24	148	1	0	0	0	0	3	0
5	0	2	0	0	0	1	177	0	2	0
6	0	1	13	157	0	1	0	0	4	7
7	1	0	0	0	164	2	0	0	0	0
8	0	99	8	7	3	0	2	2	102	2
9	0	55	2	0	5	0	1	2	6	20

```
In [7]: 1 n = 10
        2 %matplotlib inline
        3 plt.matshow(digits.images[n])
        4 print('Predict Label:', labels_kmeans[n])
        5
        6 print('Truth: ', digits.target[n])
```

Predict Label: 2

Truth: 0




## Visualization - Kmeans


```
In [8]: 1 #import libs
        2 import numpy as np
        3 from sklearn.decomposition import PCA
```

**PCA**

```
In [9]: 1 nComponents = 2
        2 vPCA = PCA(nComponents)
        3 digitData_to_2D = vPCA.fit_transform(digits.data)
        4 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_kmeans, s=
        5 plt.show())
```



```
In [10]: 1 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= digits.target, s=
        2 plt.show())
```



## Spectral clustering

```
In [11]: 1 # Spectral_clustering
2
3 from sklearn.cluster import spectral_clustering
4 from sklearn.feature_extraction import image
5 import numpy as np
6 from sklearn.neighbors import DistanceMetric
7 from sklearn.metrics.pairwise import cosine_similarity
8
9 # dist = DistanceMetric.get_metric('euclidean')
10 # graph=dist.pairwise(digits.data)
11
12 graph = cosine_similarity(digits.data)
13 label_spectral = spectral_clustering(graph, n_clusters=10)
```

## Cross table biểu thị kết quả và so sánh


```
In [12]: 1 df1 = pd.DataFrame({'labels':label_spectral,'Truth labels':digits.target})
2 ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
3 print(ct2)
```

Truth labels labels	0	1	2	3	4	5	6	7	8	9
0	1	0	0	0	163	2	0	0	0	0
1	0	0	1	147	0	0	0	0	6	2
2	177	0	1	0	1	1	0	0	0	3
3	0	0	0	4	0	157	0	0	3	3
4	0	36	115	4	0	0	0	0	1	0
5	0	0	0	16	0	20	2	0	7	133
6	0	0	2	2	11	0	0	154	3	2
7	0	2	0	0	0	2	172	0	13	0
8	0	86	53	5	5	0	7	10	101	1
9	0	58	5	5	1	0	0	15	40	36

```
In [13]: 1 n = 15
          2 plt.matshow(digits.images[n])
          3 print('lables_predict:',label_spectral[n])
          4 print(' True: ', digits.target[n])
```


lables\_predict: 3

True: 5



### Visualization - Spectral Clustering

```
In [14]: 1 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= label_spectral, s
          2 plt.show())
```




### Visualize results to compare - Using PCA

```

In [15]: 1 fig = plt.figure(figsize=(15,4))
          2 fig.suptitle('Comparison results of methods', fontsize=20)
          3
          4 ax = fig.add_subplot(1,3,1)
          5 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= label_spectral, s=
          6 ax.set_title('Spectral Clustering')
          7
          8 ax = fig.add_subplot(1,3,2)
          9 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_kmeans, s=
          10 ax.set_title('Kmeans Clustering')
          11
          12 ax = fig.add_subplot(1,3,3)
          13 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= digits.target, s=
          14 ax.set_title('Truth Result')

```

Out[15]: <matplotlib.text.Text at 0x16c59124940>



## DBSCAN

```

In [16]: 1 import numpy as np
          2 import matplotlib.pyplot as plt
          3 from sklearn.cluster import DBSCAN
          4 from sklearn import metrics
          5 from sklearn.datasets.samples_generator import make_blobs
          6 from sklearn.preprocessing import StandardScaler
          7
          8 #import scikit-learn
          9 from sklearn import metrics
          10 from sklearn.cluster import KMeans
          11 from sklearn.datasets import load_digits
          12 from sklearn.preprocessing import scale
          13 from sklearn.decomposition import PCA
          14

```

```
In [17]: 1 digits = load_digits()
2 data = digits.data
3 data = StandardScaler().fit_transform(data)
4
5 n_samples, n_features = data.shape
6 n_digits = len(np.unique(digits.target))
7 labels = digits.target
8
9 sample_size = 300
10
11 print("n_digits: %d, \t n_samples %d, \t n_features %d"
12       % (n_digits, n_samples, n_features))
13
```

n\_digits: 10,      n\_samples 1797,      n\_features 64

```
In [18]: 1 print(data)
```

```
[ [ 0.          -0.33501649 -0.04308102 ..., -1.14664746 -0.5056698
   -0.19600752]
  [ 0.          -0.33501649 -1.09493684 ...,  0.54856067 -0.5056698
   -0.19600752]
  [ 0.          -0.33501649 -1.09493684 ...,  1.56568555  1.6951369
   -0.19600752]
  ...,
  [ 0.          -0.33501649 -0.88456568 ..., -0.12952258 -0.5056698
   -0.19600752]
  [ 0.          -0.33501649 -0.67419451 ...,  0.8876023  -0.5056698
   -0.19600752]
  [ 0.          -0.33501649  1.00877481 ...,  0.8876023  -0.26113572
   -0.19600752]]
```

```
In [19]: 1 db = DBSCAN(eps=1, min_samples=1, algorithm='kd_tree').fit(data)
```

- với min\_samples = 1 => số cluster ra quá lớn = số input ban đầu

```
In [20]: 1 print(db)
```

```
DBSCAN(algorithm='kd_tree', eps=1, leaf_size=30, metric='euclidean',
        metric_params=None, min_samples=1, n_jobs=1, p=None)
```



```
In [22]: 1 import pandas as pd
2 df1 = pd.DataFrame({'labels':db.labels_, 'Truth labels':digits.target})
3 ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
4 print(ct2)
```

Truth labels	0	1	2	3	4	5	6	7	8	9
labels										
0	1	0	0	0	0	0	0	0	0	0
1	0	1	0	0	0	0	0	0	0	0
2	0	0	1	0	0	0	0	0	0	0
3	0	0	0	1	0	0	0	0	0	0
4	0	0	0	0	1	0	0	0	0	0
5	0	0	0	0	0	1	0	0	0	0
6	0	0	0	0	0	0	1	0	0	0
7	0	0	0	0	0	0	0	1	0	0
8	0	0	0	0	0	0	0	0	1	0
9	0	0	0	0	0	0	0	0	0	1
10	1	0	0	0	0	0	0	0	0	0
11	0	1	0	0	0	0	0	0	0	0
12	0	0	1	0	0	0	0	0	0	0
13	0	0	0	1	0	0	0	0	0	0
14	0	0	0	0	1	0	0	0	0	0
15	0	0	0	0	0	1	0	0	0	0
16	0	0	0	0	0	0	1	0	0	0
17	0	0	0	0	0	0	0	1	0	0
18	0	0	0	0	0	0	0	0	1	0
19	0	0	0	0	0	0	0	0	0	1
20	1	0	0	0	0	0	0	0	0	0
21	0	1	0	0	0	0	0	0	0	0
22	0	0	1	0	0	0	0	0	0	0
23	0	0	0	1	0	0	0	0	0	0
24	0	0	0	0	1	0	0	0	0	0
25	0	0	0	0	0	1	0	0	0	0
26	0	0	0	0	0	0	1	0	0	0
27	0	0	0	0	0	0	0	1	0	0
28	0	0	0	0	0	0	0	0	1	0
29	0	0	0	0	0	0	0	0	0	1
...	...	...	...	...	...	...	...	...	...	...
1766	0	0	0	0	1	0	0	0	0	0
1767	1	0	0	0	0	0	0	0	0	0
1768	0	0	0	0	0	1	0	0	0	0
1769	0	0	0	1	0	0	0	0	0	0
1770	0	0	0	0	0	0	1	0	0	0
1771	0	0	0	0	0	0	0	0	0	1
1772	0	0	0	0	0	0	1	0	0	0
1773	0	1	0	0	0	0	0	0	0	0
1774	0	0	0	0	0	0	0	1	0	0
1775	0	0	0	0	0	1	0	0	0	0
1776	0	0	0	0	1	0	0	0	0	0
1777	0	0	0	0	1	0	0	0	0	0
1778	0	0	0	0	0	0	0	1	0	0
1779	0	0	1	0	0	0	0	0	0	0
1780	0	0	0	0	0	0	0	0	1	0
1781	0	0	1	0	0	0	0	0	0	0
1782	0	0	1	0	0	0	0	0	0	0
1783	0	0	0	0	0	1	0	0	0	0

```

1784      0  0  0  0  0  0  0  1  0  0
1785      0  0  0  0  0  0  0  0  0  1
1786      0  0  0  0  0  1  0  0  0  0
1787      0  0  0  0  1  0  0  0  0  0
1788      0  0  0  0  0  0  0  0  1  0
1789      0  0  0  0  0  0  0  0  1  0
1790      0  0  0  0  1  0  0  0  0  0
1791      0  0  0  0  0  0  0  0  0  1
1792      1  0  0  0  0  0  0  0  0  0
1793      0  0  0  0  0  0  0  0  1  0
1794      0  0  0  0  0  0  0  0  0  1
1795      0  0  0  0  0  0  0  0  1  0


```

[1796 rows x 10 columns]

```

In [23]: 1 nComponents = 2
          2 vPCA = PCA(nComponents)
          3 digitData_to_2D = vPCA.fit_transform(digits.data)
          4 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= db.labels_, s=20)
          5 plt.show()

```



## Agglomerative Clustering

```

In [24]: 1 from sklearn.cluster import AgglomerativeClustering

```

```

In [25]: 1 Agglomerative_model = AgglomerativeClustering(n_clusters = 10)

```

```

In [26]: 1 db = Agglomerative_model.fit(data)

```

```

In [27]: 1 print(db.labels_)


```

[5 1 1 ..., 1 1 1]

```
In [28]: 1 import pandas as pd
2 df1 = pd.DataFrame({'labels':db.labels_, 'Truth labels':digits.target})
3 ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
4 print(ct2)
```

Truth labels labels	0	1	2	3	4	5	6	7	8	9
0	0	1	0	0	1	168	0	1	1	3
1	0	150	15	11	4	0	1	1	168	38
2	0	0	1	0	1	0	0	0	0	0
3	0	27	160	4	0	1	0	0	3	0
4	0	0	1	168	0	12	0	1	2	135
5	178	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	1	180	0	0	1
7	0	0	0	0	12	0	0	25	0	3
8	0	4	0	0	163	0	0	0	0	0
9	0	0	0	0	0	0	0	151	0	0

```
In [29]: 1 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= db.labels_, s=20)
2 plt.show()
```



## Bài thực hành 2

### Comparison of cluster methods


```
In [1]: 1 #import libs
2 from time import time
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import pandas as pd
```

```
In [2]: 1 #import scikit-learn
        2 from sklearn import metrics
        3 from sklearn.cluster import KMeans, spectral_clustering, DBSCAN, AgglomerativeClustering
        4 from sklearn.datasets import load_digits
        5 from sklearn.neighbors import DistanceMetric
        6 from sklearn.metrics.pairwise import cosine_similarity
        7 from sklearn.preprocessing import StandardScaler
```

```
In [3]: 1 digits = load_digits();
        2 print(digits.data.shape);
```

(1797, 64)

```
In [4]: 1 %matplotlib inline
        2 #plt.gray();
        3 plt.matshow(digits.images[0]);
```



## Clustering

```
In [5]: 1 #Kmeans
2 nClusters = 10
3 t0 = time()
4 kmeans_model = KMeans(nClusters)
5 labels_kmeans = kmeans_model.fit_predict(digits.data)
6 t_kmeans = time()- t0
7 #Cross table
8 print('Kmeans:\n')
9 df1 = pd.DataFrame({'labels':labels_kmeans,'Truth labels':digits.target})
10 ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
11 print(ct2)
```

Kmeans:

Truth labels	0	1	2	3	4	5	6	7	8	9
labels										
0	177	0	1	0	0	0	1	0	0	0
1	0	24	148	1	0	0	0	0	3	0
2	0	55	2	0	7	0	1	2	6	20
3	0	0	3	7	7	0	0	175	4	7
4	0	1	0	2	0	136	0	0	4	6
5	0	0	2	9	0	42	0	0	49	139
6	0	1	13	157	0	1	0	0	3	6
7	1	0	0	0	163	2	0	0	0	0
8	0	2	0	0	0	1	177	0	2	0
9	0	99	8	7	4	0	2	2	103	2

```
In [6]: 1 #Spectral_clustering
2 t0 = time()
3 graph = cosine_similarity(digits.data)
4 labels_spectral = spectral_clustering(graph, n_clusters=10)
5 t_spectral = time()- t0
6 #Cross table
7 print('Spectral clustering:\n')
8 df1 = pd.DataFrame({'labels':labels_spectral,'Truth labels':digits.target})
9 ct2 = pd.crosstab(df1['labels'],df1['Truth labels'])
10 print(ct2)
11
```

Spectral clustering:

Truth labels	0	1	2	3	4	5	6	7	8	9
labels										
0	0	58	5	5	1	0	0	15	40	36
1	0	0	0	16	0	20	2	0	7	133
2	0	2	0	1	0	2	172	0	13	0
3	177	0	1	0	1	1	0	0	0	3
4	0	36	115	4	0	0	0	0	1	0
5	0	0	2	2	11	0	0	154	3	2
6	1	0	0	0	163	2	0	0	0	0
7	0	86	53	5	5	0	7	10	101	1
8	0	0	0	4	0	157	0	0	3	3
9	0	0	1	146	0	0	0	0	6	2

```
In [7]: 1 #DBSCAN
2 data = digits.data
3 t0 = time()
4 # data = StandardScaler().fit_transform(data)
5 labels_dbscan = DBSCAN(eps=0.06, min_samples=10, metric = 'cosine').fit_predict(data)
6 t_dbscan = time() - t0
7 #DBSCAN coss table
8 print('DBSCAN:\n')
9 df1 = pd.DataFrame({'labels': labels_dbscan, 'Truth labels': digits.target})
10 ct2 = pd.crosstab(df1['labels'], df1['Truth labels'])
11 print(ct2)
```

DBSCAN:

Truth labels labels	0	1	2	3	4	5	6	7	8	9
-1	7	13	41	48	34	65	5	48	77	77
0	171	0	0	0	0	0	0	0	0	0
1	0	143	0	0	0	0	1	0	97	1
2	0	0	0	0	0	0	175	0	0	0
3	0	0	0	135	0	1	0	0	0	102
4	0	0	0	0	147	0	0	0	0	0
5	0	0	136	0	0	0	0	0	0	0
6	0	0	0	0	0	116	0	0	0	0
7	0	0	0	0	0	0	0	131	0	0
8	0	26	0	0	0	0	0	0	0	0

```
In [8]: 1 #Agglomerative Clustering
2 t0 = time()
3 Agglomerative_model = AgglomerativeClustering(n_clusters = nClusters)
4 labels_AgglomerativeClustering = Agglomerative_model.fit_predict(data)
5 t_agg = time() - t0
6 #Agglomerative Clustering - Crosstable
7 print('Agglomerative Clustering:\n')
8 df1 = pd.DataFrame({'labels': labels_AgglomerativeClustering, 'Truth labels': digits.target})
9 ct2 = pd.crosstab(df1['labels'], df1['Truth labels'])
10 print(ct2)
```

Agglomerative Clustering:

Truth labels labels	0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	179	0	0	0	2
1	0	0	0	169	0	2	0	0	1	145
2	0	27	166	0	0	0	0	0	4	0
3	0	0	1	1	3	0	0	179	1	11
4	0	0	10	13	0	0	1	0	165	2
5	0	0	0	0	178	0	0	0	0	0
6	0	0	0	0	0	1	180	0	0	0
7	178	0	0	0	0	0	0	0	0	0
8	0	59	0	0	0	0	0	0	1	20
9	0	96	0	0	0	0	0	0	2	0

## Comparison

```
In [9]: 1 n_samples, n_features = data.shape
2 n_digits = len(np.unique(digits.target))
3 sample_size = 300
4 #print frame
5 print("n_digits: %d, \t n_samples %d, \t n_features %d"
6       % (n_digits, n_samples, n_features))
7
8 print(82 * '_')
9 print('init\t\ttime\tthomo\tcompl\tv-meas\tARI\tAMI\tsilhouette')
10
11
12 data = digits.data
13 #define a function to measure and print out
14 def bench_clustering(method_name, time_, labels):
15     print('%-9s\t%.2fs\t%.3f\t%.3f\t%.3f\t%.3f\t%.3f\t%.3f'
16           % (method_name, time_,
17               metrics.homogeneity_score(digits.target, labels),
18               metrics.completeness_score(digits.target, labels),
19               metrics.v_measure_score(digits.target, labels),
20               metrics.adjusted_rand_score(digits.target, labels),
21               metrics.adjusted_mutual_info_score(digits.target, labels),
22               metrics.silhouette_score(data, labels,
23                                         metric='euclidean',
24                                         sample_size=sample_size)))
25
26
27 #Kmeans
28 bench_clustering('K-means', t_kmeans, labels_kmeans)
29 #Spectral_clustering
30 bench_clustering('spectral', t_spectral, labels_spectral)
31
32 #Agglomerative clustering
33 bench_clustering('Agg.', t_agg, labels_AgglomerativeClustering)
34 #DBSCAN
35 bench_clustering('DBSCAN', t_dbscan, labels_dbscan)
```

n\_digits: 10,      n\_samples 1797,      n\_features 64

---

init	time	homo	compl	v-meas	ARI	AMI	silhouette
K-means	0.17s	0.741	0.749	0.745	0.672	0.738	0.165
spectral	0.45s	0.711	0.716	0.713	0.624	0.708	0.170
Agg.	0.13s	0.858	0.879	0.868	0.794	0.856	0.195
DBSCAN	0.04s	0.709	0.757	0.732	0.520	0.706	0.130

## Nhận xét

- Dựa trên bảng so sánh trên, thông qua các độ đo, ta có thể thấy phương pháp Agglomerative clustering cho kết quả có độ chính xác cao nhất, với tốc độ nhanh nhất trong các phương pháp đã sử dụng

In [ ]: 1


## Visualization



```
In [10]: 1 from sklearn.decomposition import PCA
2
3 nComponents = 2
4 vPCA = PCA(nComponents)
5 digitData_to_2D = vPCA.fit_transform(digits.data)
6
7 fig = plt.figure(figsize=(15,15))
8 fig.suptitle('Comparison results of methods', fontsize=20)
9
10 ax = fig.add_subplot(3,2,1)
11 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_spectral,
12 ax.set_title('Spectral Clustering')
13
14 ax = fig.add_subplot(3,2,2)
15 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_dbscan, s=
16 ax.set_title('DBSCAN Clustering')
17
18 ax = fig.add_subplot(3,2,3)
19 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_kmeans, s=
20 ax.set_title('K-means')
21
22 ax = fig.add_subplot(3,2,4)
23 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= labels_Agglomerat
24 ax.set_title('Agglomerative clustering')
25
26 ax = fig.add_subplot(3,2,5)
27 plt.scatter(digitData_to_2D[:,0], digitData_to_2D[:,1], c= digits.target, s=
28 ax.set_title('Target Result')
29
```

Out[10]: <matplotlib.text.Text at 0x21792954828>

## Comparison results of methods




**Show centroids**

```

In [12]: 1 # create a fig to show image
          2 fig = plt.figure(figsize=(8,3))
          3 plt.title('Cluster Center visulization')
          4 plt.axis('off')
          5 # for all 0-9 labels
          6 for i in range(10):
          7     # initialize subplots in a grid 2x5 at i+1th position
          8     ax = fig.add_subplot(2, 5, 1+i)
          9
         10     # display image
         11     ax.imshow(kmeans_model.cluster_centers_[i].reshape((8,8)), cmap=plt.cm.bi
         12     #don't show the axes
         13     plt.axis('off')
         14
         15 plt.show()

```



- ( Các centroid có biểu thị hình ảnh là các chữ số )

```

In [ ]: 1

```